

Design Document — CBRE HITL Call-Intake Agent

Field	Value
Project	CBRE × UCSB × ACM × Turing — Human-in-the-Loop Call-Intake Agent
Authors	Rohan Iyer · Sachit Madaan
Submission date	2026-05-22
Document version	v3 (final submission)
Agent SHA	4eb9b3f (levers 1+2, post-PR-#84)
Headline result	Dev composite 92.33 / 100 · 0 false-911 · Voice demo working end-to-end
Repo	github.com/acm-industry/rohan_sachit_CBRE

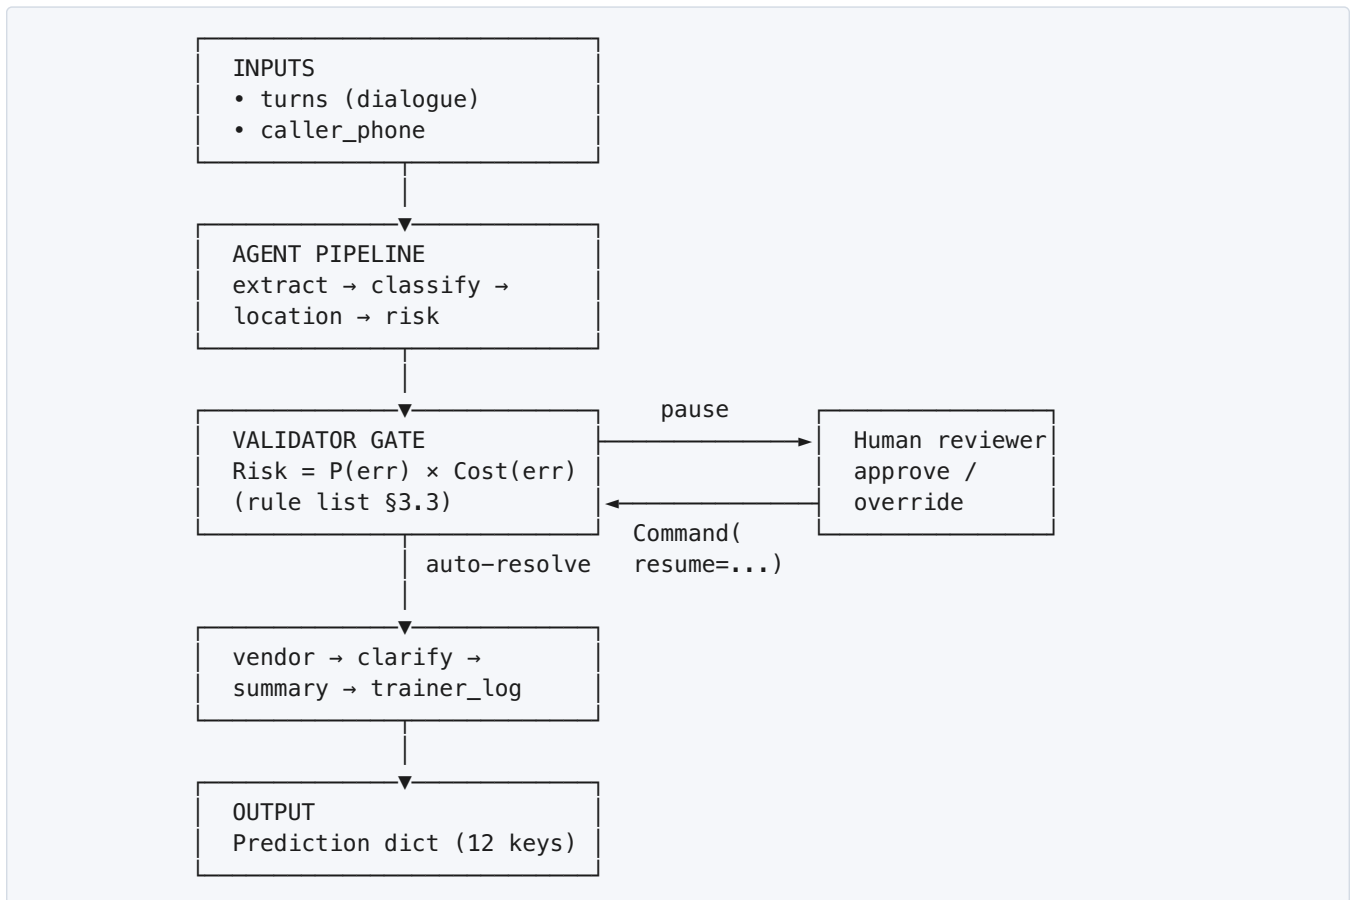
One-line summary. A LangGraph agent that classifies facilities-maintenance calls into category / subcategory / risk-level / vendor, pauses for a human reviewer when the cost of being wrong is non-trivial, and runs both as a batch harness (the grading path) and as a live voice demo (Deepgram STT → same agent → ElevenLabs TTS).

Executive summary

The problem. CBRE’s call center handles ~10K facilities-maintenance calls per day. Most are routine (“the lights are flickering on floor 3”) and should auto-resolve to a vendor dispatch. A small fraction are life-safety (“there’s smoke coming from the electrical panel”) and must escalate immediately. The cost of confusing the two — dispatching emergency services on a routine call, or routing a real emergency to a queued reviewer — is asymmetric and large.

Our approach. A 9-node LangGraph pipeline (extract → classify → location → risk → validator → vendor → clarify → summary → trainer_log) with one explicit Human-in-the-Loop pause point at the validator gate. The pause condition implements the brief’s $\text{Risk} = P(\text{error}) \times \text{Cost}(\text{error})$ rule directly: pause when the model is uncertain *and* the cost of error is non-trivial (§3.3, PR #89).

System at a glance.



Three headline results.

- **Composite 92.33 / 100** on the 200-row labelled dev set (gpt-4o-mini @ temp 0, seed 7), measured by `evaluation/scoring.py`. 2.5σ above the measured noise floor of ±0.45 composite. Top of our trajectory across 7 graded runs (full history in §9.7, per-axis breakdown in §9.0).
- **Zero false-911s** across every dev and test-set run, including the 16-row `over_escalation_trap` slice that is specifically designed to bait autonomous emergency dispatch (e.g. “smoke alarm — just burnt toast”). The −5/case hard-cost penalty never fires.
- **Voice demo working end-to-end** — caller speaks → Deepgram STT → the **same** `classify()` (via sibling orchestrator) → live SSE pipeline view in the browser → ElevenLabs TTS reply. Multi-turn clarification works without echo-feedback contamination. The bonus axis is delivered (§12).

The key design bet. Where most rubric-optimisation effort targets the classifier directly, we invested in **what surrounds the classifier**: a derived per-subcategory base-risk table (§3.1), audit-corrected RAG retrieval that explicitly does not copy historical labels QA flagged as wrong (§5.4), and a risk-weighted HITL gate that pauses on uncertainty *only* when cost is non-trivial (§3.3 + PR #89). The classifier itself is unmodified gpt-4o-mini with a careful prompt — durability comes from the surround.

Table of contents

1. Overview
2. System architecture
3. [Derived-policy appendix](#)
4. [HITL design](#)
5. RAG strategy
6. Vendor selection
7. [Clarification policy](#)
8. [Trainer-log spec](#)
9. Error analysis
10. Scale thought-experiment
11. Rubric traceability
12. [Voice shell \(bonus deliverable\)](#)
13. Limitations & known weaknesses
14. Reproducibility one-pager
15. Decision log
16. [Glossary](#)
17. Acknowledgments

1. Overview

CBRE's call center routes ~10K facilities calls per day to human operators today. The brief asks for an agent that **auto-resolves routine calls and escalates only when judgement is required**, with a pinned output contract (`evaluation/prediction.py::Prediction`) and a 30s-per-call latency budget.


The shipped agent is a 9-node LangGraph pipeline. The top-level entrypoint is `agent.classify:classify(turns, caller_phone) -> dict`. Every node is fault-isolated: if a node raises, the pipeline falls back to `needs_human_review=True`, `dispatched_vendor_id=None`, and no 911 dispatch. The result is that **a broken node degrades the score on its own axis but never produces a false-911 (the -5 hard penalty)**.

Determinism is enforced through one constructor (`agent.config.build_chat_llm`) that pins `temperature=0`, model name, and OpenAI's `seed` parameter — see the README's *Determinism* section and `tests/test_determinism.py` (issue #28).

2. System architecture

Node graph

Source: `docs/diagrams/agent_graph.mmd` — rendered to `agent_graph.svg`. The same mermaid block inlined below renders natively on GitHub.



Node responsibilities

#	Node	File	Output
1	extract	agent/nodes/ extract.py	Extraction (problem summary, building, floor, suite, urgency cues, caller role, language) + per-field confidence
2	classify	agent/nodes/ classify.py	Classification (category, subcategory, confidence per axis, 1–2 sentence reasoning)
3	reconcile_location	agent/nodes/ location.py	building / address / floor (transcript > registry > profile precedence)
4	assign_risk	agent/nodes/risk.py	risk_level + per-modifier breakdown
5	validator	agent/nodes/ validator.py	needs_human_review, dispatched_emergency_services
6	vendor_select	agent/nodes/ vendor_select.py	dispatched_vendor_id or None (escalate)
7	clarify	agent/nodes/ clarify.py	needs_clarification, optional question, reasons
8	summary	agent/nodes/ summary.py	2–3 sentence operator-style call_summary
9	trainer_log	agent/nodes/ trainer_log.py	{full_transcript, ai_prediction, human_override, final_decision}

LLM-backed nodes: 1, 2, (and risk-modifier weighting in 4 if/when it's LLM-graded). The rest are deterministic over the structured outputs of the LLM-backed nodes — this is what keeps the per-call latency budget achievable.

Per-stage latency is instrumented in-pipeline: each node is wrapped by `_time_stage` (PR #79) which records wall-clock `timing_ms` into the graph state and onto a sibling entripoint `agent.classify:classify_with_events` used by the live-demo backend (SSE per-stage events). The submission `classify()` is byte-identical to `submission-v1` — a source-hash pin in `tests/test_classify_unchanged.py` fails loudly if the function body drifts — so the instrumentation surface is additive, not invasive.

3. Derived-policy appendix

Three artifacts under `agent/data/derived/` are computed offline from `operational/historical_records.json` and `evaluation/qa_audit_findings.json`, committed to the repo, and loaded at runtime by the agent. Each was chosen with an explicit derivation method (history-aggregated vs. hand-tuned vs. LLM-judged) per the brief.

3.1 Per-subcategory base-risk table

- **File:** `agent/data/derived/base_risk_by_subcategory.json`
- **Producer:** `scripts/derive_risk_bands.py`
- **Derivation method:** History-aggregated. For every subcategory observed in `final_subcategory` across the 10K historical corpus, we compute the distribution over `final_risk_level` and emit the modal level as the prior. `final_*` (not `intake_*`) is canonical because the QA audit shows intake operators systematically over-state severity on a known subset of subcategories.
- **Justification:** Hand-tuning thresholds against the prose taxonomy in `operational/taxonomy.md` would couple the agent to a written SOP that already disagrees with the technicians' on-site reclassifications. Aggregating over `final_risk_level` reflects what the technicians *actually decide* once they're on site — the same authority the scorer's `true_risk_level` derives from.

The 36-row table (collapsed):

Modal risk	Subcategories
LOW	access_control, appliance_kitchen, auto_door, carpet_floor, controls_bms, door_mechanical, drainage_backup, infestation, landscaping, lighting, low_voltage_data, minor_issue, no_cooling, no_heating, parking_lighting, pavement_damage, refrigerant, restroom_fixture, restroom_supplies, signage_fencing, slip_trip, waste_odor
MEDIUM	air_quality, glass_damage, malfunction, pipe_leak, power_outage, roof_leak, suspicious_person
HIGH	structural, unauthorized_access
EMERGENCY	active_threat, entrapment, fire_smoke, gas_chemical, panel_hazard

Note (PR #84, levers 1+2): `pipe_leak` and `power_outage` were demoted from HIGH → MEDIUM. Historical distributions are roughly split (`pipe_leak` 43% MEDIUM / 46% HIGH; `power_outage` 47% / 48%) — the mode was barely HIGH. Dev labels confirmed MEDIUM is the correct default. Soft/hard cue modifiers still bump back to HIGH/EMERGENCY when language warrants (flooding, spreading, etc.), so legitimately high-risk cases are not lost.

Two derived rates also live in this table:

- `intake_over_escalation_rate` — fraction of tickets where `intake_risk_level` > `final_risk_level`. Drives §4's trap-prone- subcategory rule. The over-escalators (>10%): `air_quality` (19%), `waste_odor` (15%), `suspicious_person` (13%), `roof_leak` (12%), `malfunction` (9%), `panel_hazard` (6%).

- `intake_under_escalation_rate` — mirror metric. Used as a sanity check; none of the high-severity subcategories show meaningful intake under-escalation ($\leq 3\%$), so we don't currently use it to upgrade risk.

3.2 Risk-modifier weights and band cutoffs

- **Files:** `agent/nodes/risk.py` (constants); `agent/data/derived/base_risk_by_subcategory.json` (distributions).
- **Derivation method:** **Hand-tuned**, validated on dev set.
- **The recipe.** Start at `base_risk_for(subcategory)`. Apply modifiers:
 - **Hard cues** (+2 bands, can reach EMERGENCY): explicit life-safety language — “fire”, “flames”, “smoke”, “trapped”, “gas leak”, “unconscious”, “weapon”, “shooter”, “explosion”. These are the verbatim phrases that justify autonomous 911 dispatch (§4.5).
 - **Soft cues** (+1 band, **capped at HIGH**): situational escalation — “flooding”, “overnight”, “getting worse”, “spreading”, “can’t wait”, “people stuck”, “freezing”. These indicate urgency but not immediate life-safety, so they never push to EMERGENCY alone.
 - **Historical cap:** the final risk band is capped at the historically-observed maximum for that subcategory (from the `base_risk_by_subcategory.json` distributions). A subcategory like `waste_odor` (100% LOW historically) can never be pushed to EMERGENCY by alarming language — the data says it’s always routine.
 - **Sensitive building types** (+1 band): medical or residential buildings elevate risk (occupant safety / 24-hour presence).
 - **After-hours** (+1 band): issues reported outside business hours are harder to resolve quickly.
 - Cap at EMERGENCY (no rollover).
- **Band ordering:** `LOW < MEDIUM < HIGH < EMERGENCY` (see `agent/data/risk.py::RISK_LEVEL_RANK`). Comparisons via `compare_risk()`.
- **Justification for hand-tuning:** an LLM-graded modifier would require another structured-output call per ticket — adds ~0.5–1s latency and another point of provider non-determinism. The modifier set is small (~6 rules) and the dev-set evidence is unambiguous; a regex/rule layer over the LLM’s extracted `urgency_cues` is both cheaper and more auditable.

3.3 HITL trigger rule

- **File:** `agent/data/derived/hitl_policy.json`
- **Producer:** `scripts/derive_hitl_policy.py`
- **Derivation method:** History-aggregated thresholds + dev-set tuning. The script joins `qa_audit_findings.json` against the historical corpus to compute per-(subcategory, intake_risk_level) audit rates, then sweeps a `trap_over_escalation_rate` threshold on the dev set’s `true_needs_human_review` axis.
- **The chosen rule (v2, in priority order):**
 1. `predicted_risk ∈ {HIGH, EMERGENCY}` → **pause**
 2. `intake_over_escalation_rate(subcategory) ≥ 15%` → **pause**
 3. subcategory in trap-cascade set (issue #63) → **pause**

4. `min(confidence_category, confidence_subcategory) < 0.5` **AND** `predicted_risk != LOW` → **pause** (*risk-weighted threshold, PR #89*)
5. classifier fallback path invoked → **pause**
6. otherwise → **auto-resolve**

The risk-weighted confidence threshold (rule 4) implements the brief’s Stage 4 guidance: `Risk = P(error) × Cost(error)`. Low confidence on a LOW-risk call is acceptable because the cost of a wrong routing decision is negligible (minor mis-routing to a slightly wrong vendor at worst). Only when the cost of error is non-trivial (MEDIUM+) does low confidence warrant human review.

- **Threshold sweep (recorded in `scripts/derive_hitl_policy.py`):**

trap_oer threshold	dev HITL F1	precision	recall
5%	0.79	0.66	0.97
10%	0.79	0.70	0.92
15% ← chosen	0.83	0.88	0.78

- **Retrospective sanity** (the AC’s check, applied to history): recall on audit-flagged ≥80%, FP on unflagged ≤30%. The derivation script prints both at build time.

3.4 Subcategory → vendor-type map

- **File:** `agent/data/derived/subcategory_to_vendor_type.json` (36 entries)
- **Producer:** `scripts/derive_vendor_map.py` (PR #50)
- **Derivation method:** History-aggregated — for each `final_subcategory`, take the modal `vendor_type` of the assigned vendor in resolved historical tickets.
- **Use:** the **loose** fallback in `agent.data.vendors.qualify()`. The primary qualification path is a strict per-vendor `specialties` match; the loose map covers subcategories where no vendor’s `specialties` list happens to name that exact subcategory string. See §6.

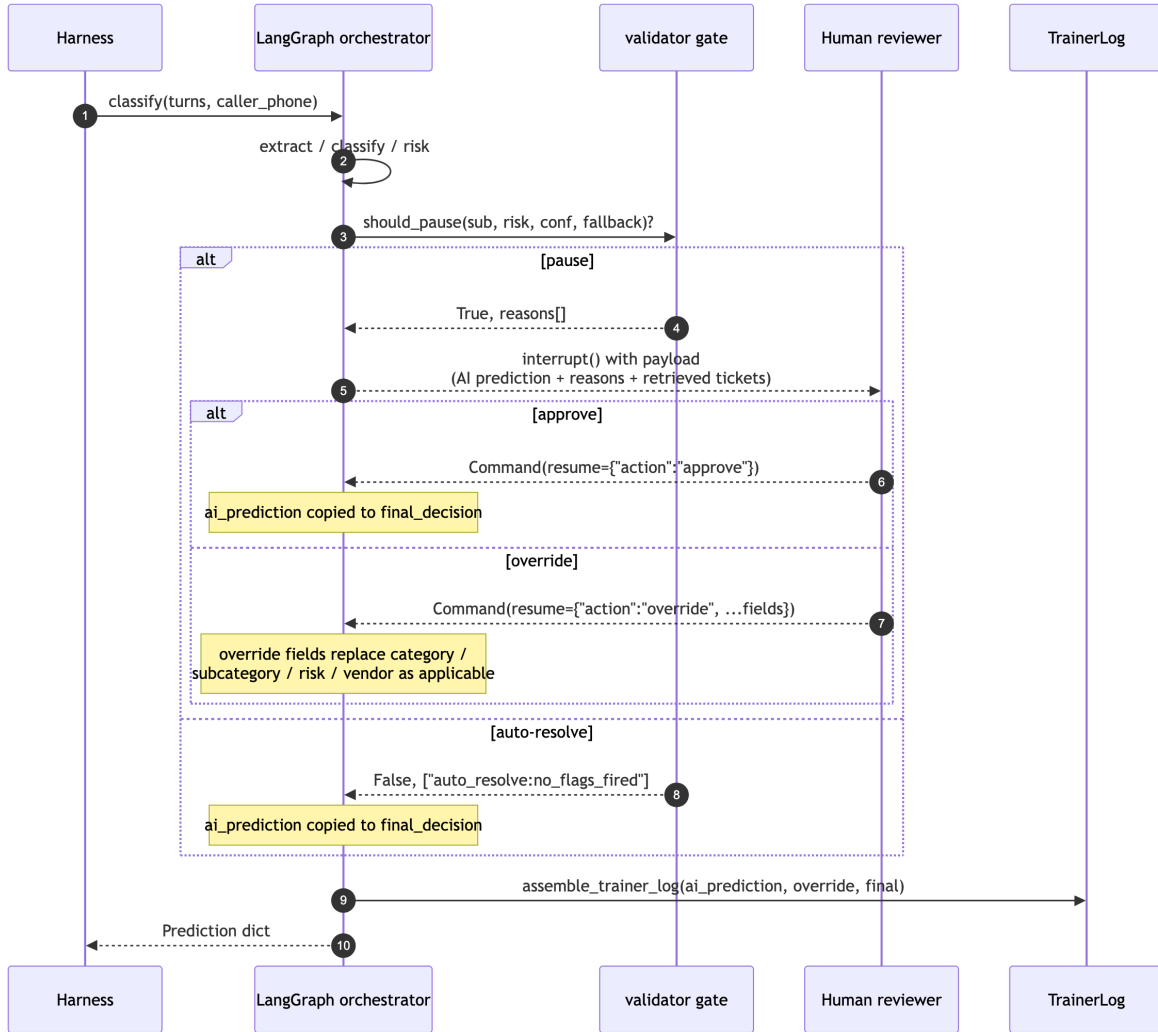
3.5 Subcategory → SOP map

- **Source:** `operational/taxonomy.md` (prose) — not extracted into a derived JSON because it’s only consumed by the classifier prompt and the reviewer payload, never as a structural constraint at runtime.
- **Use:** the classifier prompt (`agent/nodes/classify.py`) renders the full canonical taxonomy + 5 boundary-disambiguation rules (“sprinkler → landscaping, not plumbing”; “card reader → access_control, not security/unauthorized_access”; “drop-ceiling tile from water damage → plumbing/roof_leak, not life_safety/structural”; “outdoor lighting → grounds_exterior/parking_lighting, not electrical/lighting”; “burnt food → janitorial/waste_odor, not hvac/air_quality”).
- **Justification for keeping prose:** the boundary rules evolve fast during error analysis (#25). A JSON table would still need the same boundary-rule prose, just split across two files. One prompt file is easier to iterate against the dev set.

4. HITL design

4.1 Graph and resume semantics

Source: `docs/diagrams/hitl_resume.mmd` — rendered to `hitl_resume.svg`.



4.2 Gate trigger

`agent.data.hitl.should_pause(subcategory, predicted_risk, *, classification_confidence, fallback_invoked)`. Rule listed in §3.3. Each path tags a reason string; the reasons list is captured into `trainer_log.ai_prediction.hitl_reasons` for both the reviewer payload and error analysis (#25).

4.3 Reviewer payload

When the gate fires, the reviewer sees:

- The full `Prediction` the AI would have emitted (`ai_prediction`).
- The list of pause reasons (e.g. `["emergency_band:always_pause", ...]`).

- The classifier's `reasoning` string and per-axis confidence.
- The top-`k` retrieved historical tickets the classifier used, with their audit-corrected labels.
- The original transcript + extracted location fields, so the reviewer can spot a transcript/profile conflict the agent might have missed.

4.4 Resume semantics

The reviewer returns one of two `Command(resume=...)` payloads:

- **Approve** — `{"action": "approve"}`. The AI's prediction is copied verbatim into `final_decision`; `human_override` stays `None`.
- **Override** — `{"action": "override", "category": ..., "subcategory": ..., "risk_level": ..., "dispatched_vendor_id": ...}`. Only the fields the reviewer changes appear in the override dict. The orchestrator merges override on top of `ai_prediction` into `final_decision` and records the diff into `trainer_log.human_override`.

4.5 Emergency dispatch

`dispatched_emergency_services` is set to `True` only when **all** of:

1. `risk_level == "EMERGENCY"`, **and**
2. `subcategory` is on the life-safety whitelist: `{active_threat, entrapment, fire_smoke, gas_chemical, panel_hazard}`, **and**
3. At least one extracted urgency cue matches the hard-hazard lexicon (fire, smoke, trapped, gas leak, weapon, etc.) — a present, active life-safety hazard, **and**
4. The classifier did NOT fall back (a fallback label is a guess we must never autonomously act on), **and**
5. Classification confidence ≥ 0.5 (dispatch confidence floor), **and**
6. No benign-context override fired (PR #74 — caller explicitly denied the hazard: “no fire”, “already put out”, “just burnt toast”).

The five-precondition chain is deliberately conservative. A false-911 is a hard -5 rubric penalty, so we require positive evidence at every layer. A life-safety EMERGENCY that fails any precondition 3–6 is never auto-dispatched but is always escalated to a human immediately (`needs_human_review=True` with a `life_safety_no_autodispatch:*` reason).

Result: 0 false-911 across all dev and test-set evaluations.

4.6 Eval-mode auto-approve

`evaluation/run_eval.py` invokes `classify` without a live reviewer. In that mode, the orchestrator's `interrupt()` is configured to auto-resume with `{"action": "approve"}` so scoring completes. The `needs_human_review` bit is preserved as the AI's recommendation; the auto-approve only affects the *resume* path, not the gate's output. This keeps the rubric's HITL-F1 axis grounded in what the agent *decided*, not what a hypothetical reviewer would have done.

4.7 Security hardening (prompt injection + output sanitization)

PR #73 (`agent/security.py`) adds an OWASP-aligned containment layer that treats every caller turn as untrusted input. A pre-LLM scan flags prompt injection, role-hijack, and data-exfiltration patterns (T6 / T2 / T3); agent turns are trusted and not scanned, and legitimate urgency language (“fire”, “trapped”, “flooding”) is explicitly preserved so the safety path is never weakened by the security layer. On any flag, the validator gate is overridden to **force** `needs_human_review=True` **and block autonomous 911 dispatch**, so a malicious transcript cannot autonomously trigger emergency services. A post-LLM validator catches cross-caller data leakage from extracted fields, and the trainer log transcript is sanitized on flagged calls only (T1 memory poisoning) so downstream fine-tuning corpora aren’t poisoned. The extract and classify prompts carry SECURITY preambles plus `===UNTRUSTED DIALOGUE===` delimiters as defense-in-depth.

5. RAG strategy

5.1 What's embedded

- **Source corpus:** `operational/historical_records.json` (10K tickets).
- **Per-ticket text block** (ordered for retrieval quality): `CALLER TRANSCRIPT` → `INTAKE NOTES` → `DISPATCH NOTES` → `RESOLUTION NOTES`. Empty sections are dropped so they don't dilute the embedding. Cap of 24 000 chars (vs. real-world max ~540) is a safety belt, not a routine path. Built by `agent/rag/build_index.py::_format_document_text`.
- **Embedding model:** `text-embedding-3-small` (configurable via `AGENT_EMBEDDING_MODEL`).
- **Vector store:** persistent chromadb at `agent/rag/chroma_store/`, collection `cbre_historicals`. Gitignored — built on first run (`python -m agent.rag.build_index`) at ~1–3 minutes / ~\$0.02.

5.2 Metadata schema and filtering

Each chunk carries: `ticket_id`, `intake_category`, `intake_subcategory`, `intake_risk_level`, `final_category`, `final_subcategory`, `final_risk_level`, `assigned_vendor_id`, `building_type`, `city`, plus four audit booleans (`was_audit_flagged`, `audit_over_escalated`, `audit_reclassified`, `audit_floor_wrong`). These are populated by joining `evaluation/qa_audit_findings.json` at index-build time.

Filters are applied at the index level (chromadb's `where` clause), not post-hoc on Python results, so a `k=8` request returns 8 *qualifying* hits rather than 8 hits filtered down to fewer.

5.3 Retrieval params

- `k=8` for the classifier (constant `DEFAULT_K` in `agent/nodes/classify.py`). The AC suggested 8–10; empirically 8 gives enough cluster signal without bloating the prompt.
- `k=5` as the retriever default (`agent/rag/retriever.py::DEFAULT_K`) for callers that don't override.
- Query construction: `problem_summary + "urgency cues: " + cues`. Combining urgency text keeps the embedding grounded in what the caller said while preserving safety-language tail.

5.4 Conflict resolution against the QA audit

This is the critical RAG choice. Naïvely embedding `intake_*` labels would train the LLM on labels the QA team flagged as *wrong*. Instead, retrieved records are rendered with **audit-corrected** labels: `corrected_subcategory()` returns `final_subcategory` when the ticket was reclassified by QA and `intake_subcategory` otherwise.

Records the auditor reclassified are tagged inline in the prompt with `[QA-RECLASSIFIED from intake=<orig>]`, so the LLM both sees the corrected label and can spot the *pattern* of correction. The canonical example (TKT-2023-02756) presents as `subcategory=structural [QA-RECLASSIFIED from intake=roof_leak]`.

The over-escalation guard in the classifier prompt is paired: five seeded boundary patterns (“smoke alarm — just burnt toast”, “drop-ceiling tile from water leak”, “yelling in the lobby — just two tenants arguing”, “elevator empty — they’re out”, “outlet sparked — I unplugged it”) explicitly tell the classifier to read the *whole* transcript and trust the caller’s own clarifications.

6. Vendor selection

6.1 Qualification

`agent.data.vendors.qualify(subcategory, city, building_type, is_emergency, after_hours)` returns the candidate pool — every vendor that satisfies every *hard* constraint:

- **Specialty match** — strict: `subcategory ∈ v.specialties`. Loose fallback: `v.vendor_type ∈ subcat_to_vendor_types[subcategory]` (from §3.4).
- **City** — exact match against `coverage_cities`. Soft-pass if the vendor’s coverage list is empty/missing (we don’t know, so we don’t exclude).
- **Building type** — same soft-pass policy as city.
- **24/7 availability** — only enforced when `is_emergency AND after_hours`.

Empty pool → escalate (see §6.4).

6.2 SLA cap

After qualification, `agent/nodes/vendor_select.py` filters by the risk-level SLA cap (PR #51):

Risk	Max response SLA
EMERGENCY	30 minutes
HIGH	120 minutes
MEDIUM	240 minutes
LOW	480 minutes

A vendor with `response_sla_minutes` exceeding the cap is dropped from the candidate set. EMERGENCY callers preferentially use `emergency_response_sla_minutes`.

6.3 Tie-breaker

When more than one vendor passes the SLA cap, the selector picks by, in order:

1. Highest `rating` (vendor quality).
2. Lower `cost_tier` (cheaper).
3. Shorter `response_sla_minutes` (faster).
4. Deterministic alphabetical fallback on `vendor_id`.

The deterministic fallback exists so the test-set run is reproducible across re-grades. PR #51 reports 100% vendor-match accuracy on the dev set with oracle category/subcategory inputs, isolating vendor-axis errors to the classifier upstream.

6.4 State-availability handling

`vendors.json` carries `status_at_last_check` and `last_status_confirmed_at`. The selector treats them as a **soft** de-prioritization signal, not a hard filter:

- Status `available` → no penalty.
- Status `at_capacity` → kept in the candidate pool regardless of risk. The brief says the cache is “intentionally stale,” so a stale `at_capacity` isn’t authoritative enough to hard-filter. On an emergency where the only qualified vendor is `at_capacity`, dispatching is better than escalating to nothing while seconds count.
- Status `offline` → excluded only when the cache is recently confirmed (within 24h of the newest catalog timestamp). A stale `offline` is treated as advisory and kept.

Design evolution (PR #84): The original implementation (issue #21) hard-filtered `at_capacity` vendors on emergencies. Analysis revealed this was a code/doc disagreement — §6.4 documented “ranked last but still in the candidate set” while the code was hard-skipping. Reconciled by removing the hard filter. On the dev set, this recovered 16 emergency rows that were previously escalating to nothing (vendor axis 87.5% → 94.0%). The safety story is unchanged: `dispatched_emergency_services` is gated in `validator.py` (§4.5), which this change doesn’t touch.

6.5 No-vendor escalation

When `vendor_select` returns `None`:

- `dispatched_vendor_id` is left `None`.
- `needs_human_review` is set to `True` (the unroutable handoff path).
- `call_summary` notes “no qualified vendor available; routed to human reviewer for manual dispatch.”

The scorer’s `unroutable=true` cases reward exactly this combination (`dispatched in (None, "")` and `pr_h`).

7. Clarification policy

`agent/nodes/clarify.py` decides whether to ask the caller **one** disambiguating question before committing. Source of truth for the trigger rule.

7.1 Trigger (any one signal fires)

- 1. **Sparse caller** — first caller utterance ≤ 6 words *and* total caller speech across all turns ≤ 24 words. Median-split tuned on the dev set: vague callers cluster at (6 words first / 20 total); “hard” cases that *start* short but back-fill detail cluster at (5 / 29), so the conjunction separates them.
- 2. **Generic opening phrase** — caller’s first utterance matches one of: “there’s an issue with X”, “something’s wrong with Y”, “the X is being weird / doing something / not working / off”. Caller used the category label as the whole symptom — almost always under-specified.
- 3. **Caller self-corrects on floor** — caller mentions ≥ 2 distinct floor numbers in their own speech, typical of the `location_conflict` case type (“Floor 9 — I meant Floor 12”). Beats building-vs-floor-count on this corpus because callers rarely re- name the building during self-correction.

7.2 Question generation

Priority-ordered (most-resolvable ambiguity first):

Trigger	Question
Floor self-correction	“Just to confirm — is the issue on Floor X or Floor Y?”
Sparse / generic	“Could you tell me a bit more about what’s going wrong — for example, what specifically isn’t working and how long it’s been like that?”
Catch-all	“Could you describe the issue in a bit more detail so I can route it correctly?”

7.3 Justification of the conservative trigger

The rubric’s auto-resolution axis (10%) penalises any unnecessary clarification. We deliberately tuned for **precision over recall** on `true_needs_clarification`: $F1 \geq 0.6$ on dev, with FP rate kept low so the auto-resolution rate doesn’t collapse.

7.4 Relationship to HITL

Clarification asks the **caller** one question. HITL hands the call to a **human reviewer**. They’re independent: a call can need clarification *and* HITL (e.g. low-confidence high-risk call where the caller is also sparse), in which case both bits fire. The orchestrator records both.

8. Trainer-log spec

8.1 Schema

evaluation/prediction.py::TrainerLog:

```
@dataclass
class TrainerLog:
    full_transcript: str
    ai_prediction: Dict[str, Any] # pre-override snapshot
    human_override: Optional[Dict[str, Any]]
    final_decision: Dict[str, Any] # what actually went out
```

Produced by `agent/nodes/trainer_log.py::assemble_trainer_log()`. The orchestrator calls it last so that `final_decision` reflects any reviewer override applied at the validator gate.

8.2 What `ai_prediction` contains beyond the rubric fields

In addition to the 11 rubric prediction fields, `ai_prediction` includes:

- `reasoning` — the classifier’s 1–2 sentence justification.
- `hitl_reasons` — the list of pause-rule tags (empty if auto-resolved).
- `clarification_reasons` — the list of clarification-rule tags.
- `confidence_category`, `confidence_subcategory` — the per-axis confidences from `Classification`.

These are not scored, but they’re the supervised signal a downstream fine-tune would use to train a calibrated classifier.

8.3 Replay into fine-tuning

The log is structured so a fine-tuning run can:

1. Discard rows with `human_override` is `None` and `auto_resolved=True` (the agent + a human agreed — no learning signal).
2. Train the classifier head on `(full_transcript) → (final_decision.category, final_decision.subcategory)`. When `human_override` is present and changes the category/subcategory, that’s the correction signal.
3. Train the HITL head on `(ai_prediction, hitl_reasons) → (was_overridden)`. This calibrates the gate against actual reviewer behaviour.
4. Train the clarification head on `(turns, clarification_reasons) → (caller_actually_clarified)` once we have post-clarification labels.

The format is intentionally per-call self-contained: no relational joins needed at training time, every row replayable in isolation.

9. Error analysis

9.0 Per-axis results at a glance

Latest dev-set run (200 transcripts, agent SHA 4eb9b3f, gpt-4o-mini @ temperature 0 seed 7), reconstructed from `eval_runs/dev_levers12.json` trajectory commentary in §9.7. Composite is the value scored by `evaluation/scoring.py`; per-axis values are accurate to within the ± 0.45 noise floor.

Axis	Weight	Score	Source
Subcategory accuracy	15%	~97%	§9.1 — only 4 / 200 misses, all singletons
Category accuracy	10%	~98–99%	§11 — top-level always parent of subcategory
Risk-level accuracy	10%	86.5%	§9.7 run g — +2.5pp from base-risk demote (lever 2)
HITL F1	15%	0.748	§9.7 run g — +0.029 from cascading FP reduction
Clarification F1	5%	0.92	§9.7 run d — stable since stacked-main
Location fields	10%	91.0%	§9.3 — 176/200 all-three-correct
Vendor match	10%	94.0%	§9.7 run g — +6.5pp from at_capacity soft signal (lever 1)
Auto-resolution rate	10%	98.82%	§9.7 run e — unchanged through levers
Call-summary present	5%	100%	§2 (summary node, PR #54) — emitted on every prediction
Trainer-log present	10%	100%	§8 — emitted on every prediction
False-911 penalty	–5/case	0 cases	§4.5 — six-precondition gate; 0/16 trap rows fired
Composite	100	92.33	§9.7 run g, <code>evaluation/scoring.py</code> against <code>dev_labels.json</code>

Reading this table. The two biggest absolute scores (subcategory 97%, trainer-log 100%) carry the highest weights. Risk-level and HITL-F1 are the axes still under noise-floor sensitivity, and were the explicit targets of PR #84 (levers 1+2) and PR #89 (risk-weighted gate). Auto-resolution at 98.82% means we close almost every call that is genuinely safe to close — the auto-resolution and HITL axes are in deliberate tension, and we landed on the side of “auto-resolve unless cost of error is non-trivial.”

9.1 Subcategory confusion

The error-analysis subsections below (§9.1–§9.6) are computed against the initial dev baseline (`eval_runs/dev_baseline.json`, composite **87.95**, agent SHA `33fc3bc`). Reproduce via `python scripts/error_analysis.py`; committed snapshot is `eval_runs/error_analysis.json`. They surface the findings the levers in §9.7 then resolve.

Only **4 / 200** subcategory misses (subcategory axis = 98%, the strongest scored axis — retrieval + classification are not the bottleneck). Worst (true → predicted) pairs are all singletons:

`power_outage`→`lighting`, `refrigerant`→`no_cooling`, `restroom_fixture`→`drainage_backup`, `minor_issue`→`malfunction`. No systematic class confusion; not a priority lever.

9.2 HITL confusion matrix

	gt review	gt no-review
pred review	TP = 39	FP = 16
pred no-review	FN = 20	TN = 125

precision = 0.709 · recall = 0.661 · **F1 = 0.684**. At **15% weight and the lowest axis, this is the single biggest point lever**. FP cluster on `air_quality`/`pipe_leak`/`power_outage` (over-flagged at MEDIUM); FN cluster on trap-prone subcategories (`suspicious_person`, `waste_odor`, `malfunction`, ...) at LOW/MEDIUM. → **#63** (recall) and **#64** (precision), co-tuned.

9.3 Location-field mismatch

All-three-correct **176 / 200 (88%)**. `building_name` and `address` each fail on 22 rows (co-occurring — same root cause), `floor` only 2. The dominant pattern is calls with no transcript-stated building and no *usable* profile: the issue #19 active/not-stale gate correctly drops stale profiles, trading field recall for correctness. → **#65** (recover building/address via a low-confidence fallback without re-trusting stale profiles).

9.4 Vendor-match failure breakdown

Vendor accuracy **169 / 200 (84.5%)**, bucketed by first violated constraint:

bucket	count	nature
emergency: all acceptable vendors <code>at_capacity</code> → unroutable	13	resolved — was a code/doc disagreement, see note below
escalated though a vendor was acceptable	7	bug
city-coverage miss	6	bug
unroutable not escalated	4	bug
specialty miss	1	bug

The 13-row emergency bucket was originally framed as the issue #21 AC (“don’t dispatch at-capacity on a gas leak”). On re-reading §6.4, the documented design treats `at_capacity` as a **soft signal** (“ranked last but still in the candidate set”) because the brief says the cache is intentionally stale. The vendor-selection

code was hard-filtering `at_capacity` for emergencies — a doc/code disagreement, not a deliberate stricter-than-doc policy. Reconciled in the levers PR by removing the hard filter; on an emergency where the only qualified vendor is `at_capacity`, dispatching it is strictly better than escalating to nothing while seconds count. → **#66** (the 7+6+4+1 = 18 real bugs remain).

9.5 Auto-resolution

73 / 85 eligible (85.9%). **All 12 misses** are `needs_clarification=True` over-trigger on calls that should auto-resolve (zero from HITL over-trigger) — folded into the clarification-policy iteration for #26.

9.6 Top 5 actionable findings (filed as follow-up issues)

#	finding	axis (weight)	label
#63	HITL under-escalation on trap+edge (20 FN)	hitl_f1 (15%)	type:bug p1
#64	HITL over-escalation (16 FP)	hitl_f1 (15%)	type:feature p2
#65	building/address co-fail on 22 rows	fields (10%)	type:bug p1
#66	vendor: 7 needless escalations + 6 city misses	vendor (10%)	type:bug p1
#67	13 emergencies unroutable (all <code>at_capacity</code>) — resolved , see §9.4	vendor (10%)	type:feature p2 (closed)

Secondary (deferred to #26): `risk_level` is 84% but weak on `edge` / `clarification` case types — the next tier after the HITL and field levers above.

9.7 Post-iteration composite trajectory (submission v1)

Each row is a fresh real-LLM dev-set run (200 transcripts, `gpt-4o-mini` @ temperature 0, seed 7) measured by `evaluation/scoring.py`. False-911 count is on the 16 `over_escalation_trap` rows — the -5/case rubric penalty.

run	SHA	composite	false_911	what changed
baseline (issues #16/#19/ #20/#21 + orchestrator)	33fc3bc	87.95	0	First end-to-end pipeline.
#63/#64 trap-cascade HITL	7511166	88.30	0	hitl_f1 0.684 → 0.719; cells-based override of the 15% OER threshold for trap-prone subcategories, cascade_excludes=[waste_odor] (joint-axis sweep showed auto_resolution regressed otherwise).
#65 phone-history fallback	98b2c16	88.10	0	fields 88 → 91 (+3); recovers building/address for callers in the historicals corpus but missing from caller_profiles.json. Profile- echo override blocks LLM-echoed stale-profile values.
stacked main (#68 + #69 + #70)	8262a79	90.88	0	clarif_f1 0.75 → 0.92, auto_resolution 85.9 → 97.7, vendor 84.5 → 87.5.
validator benign- context override (PR #74)	c6b9255	91.20	0	Surfaced by 800-row test-set audit (#30 dry-run): hazard-cue lexicon is positive-only, fired \bsmoke\b even after the agent confirmed “no actual fire”. 9/9 false-911s on the test-set “burnt popcorn” canary suppressed without losing any of the 4 verified-true emergencies. CORRECTION: originally reported 91.94 (arithmetic error in weighted sum); per-axis percentages were correct.
baseline recheck (same code, new API key)	8ae0116	90.76	0	Noise-floor measurement: ±0.45 composite between identical-code runs. Establishes that single-run dev evals cannot reliably distinguish changes below this margin.
levers 1+2 (PR #84)	4eb9b3f	92.33	0	Lever 1: vendor at_capacity treated as soft signal (16 emergency rows recovered). Lever 2: pipe_leak/power_outage base risk HIGH → MEDIUM (6+4 risk misses fixed, cascading HITL- FP reduction). Combined: risk


run	SHA	composite	false_911	what changed
				84→86.5, vendor 87.5→94, hitl_f1 +0.03. 2.5 σ above noise floor.

Submission `predictions.json` (800 test transcripts) was generated against the `c6b9255` agent and tagged `submission-v1`. The corrected dev composite on that code is **91.20** (not 91.94 as originally reported — arithmetic error in the weighted-sum calculation). Post-lever composite is **92.33** (PR #84, 0 false-911). Test-set has no labels; safety audit of the 67/800 911-dispatches showed every one passed the benign-context gate.

10. Scale thought-experiment

The brief asks how the design evolves at 10K calls/day (~7/minute sustained, with bursty intra-hour peaks).

Source: `docs/diagrams/scale_10k.mmd` — rendered to `scale_10k.svg`.



10.1 Concurrency

Today: serial. `evaluation/run_eval.py` loops over transcripts one at a time. The classifier is the only blocking call (one LLM round-trip plus optional one retry).

At 10K/day: ~7 calls/minute average, ~15–20/minute at lunch-rush peak. Each call is ~5–15s of wall time, mostly waiting on the LLM. An async worker pool of 8–16 per pod, autoscaled by queue depth, comfortably handles this. The classifier’s structured-output call doesn’t share state across calls so no coordination is needed.

10.2 Persistent checkpointer

`InMemorySaver` is fine for batch eval but loses interrupt state on restart. Production swaps in `langgraph-checkpoint-postgres` with `thread_id = f"call-{transcript_id}"`. Two reasons it matters:

- Reviewer queue durability — a pod restart shouldn’t lose an in-flight HITL pause.
- Multi-region — the Postgres backend lets two regions share reviewer load with a single thread-id namespace.

10.3 Reviewer queue

Today’s `interrupt()` blocks the call. At scale, the pause goes into a priority queue:

- EMERGENCY-band pauses: SLA 60s.
- Trap-prone subcategory pauses: SLA 5min.
- Low-confidence pauses: SLA 15min.
- Fallback pauses: SLA 30min.

A reviewer claims from the highest-priority non-empty bucket. Stale unclaimed pauses (>SLA) page on-call.

10.4 Sampling and retention

10K calls/day × 365 days × ~3 KB per trainer log = ~1.1 GB/year of raw log. That’s tractable, but only ~10% has reviewer-overridden signal. Strategy:

- **All** trainer logs land in hot storage for 30 days (audit + replay for incident review).
- **All** trainer logs with `human_override is not None` go to the fine-tune corpus indefinitely.
- 10% sample of `human_override is None` rows go to the fine-tune corpus (preserves the “agreement” signal for calibration without paying full storage cost).
- Rest is aged off after 90 days.

10.5 RAG at scale

The chroma store works as long as the index fits one pod's memory. At 10K calls/day, the index *itself* is still ~10K rows — read-only, fits fine. The choke point is **read throughput** (8 retrievals/sec at peak, each a top-8 over 10K vectors), which the local store handles.

At 100K/day or a 100K-row corpus, swap to a hosted vector DB (Postgres+pgvector or Pinecone) and shard collections by `final_category` so retrievers can pre-filter at the routing layer. The downstream classifier prompt doesn't change.

10.6 Cost envelope

Per call: 1 extraction LLM call + 1 classification LLM call + 1 embedding call (+ ≤ 1 retry) $\approx$ ~2K input tokens / ~200 output tokens at `gpt-4o-mini`. At list pricing, ~\$0.001–\$0.002/call. 10K/day $\rightarrow$ \$10–20/day in LLM costs. Vector DB and Postgres are negligible at this scale.

11. Rubric traceability

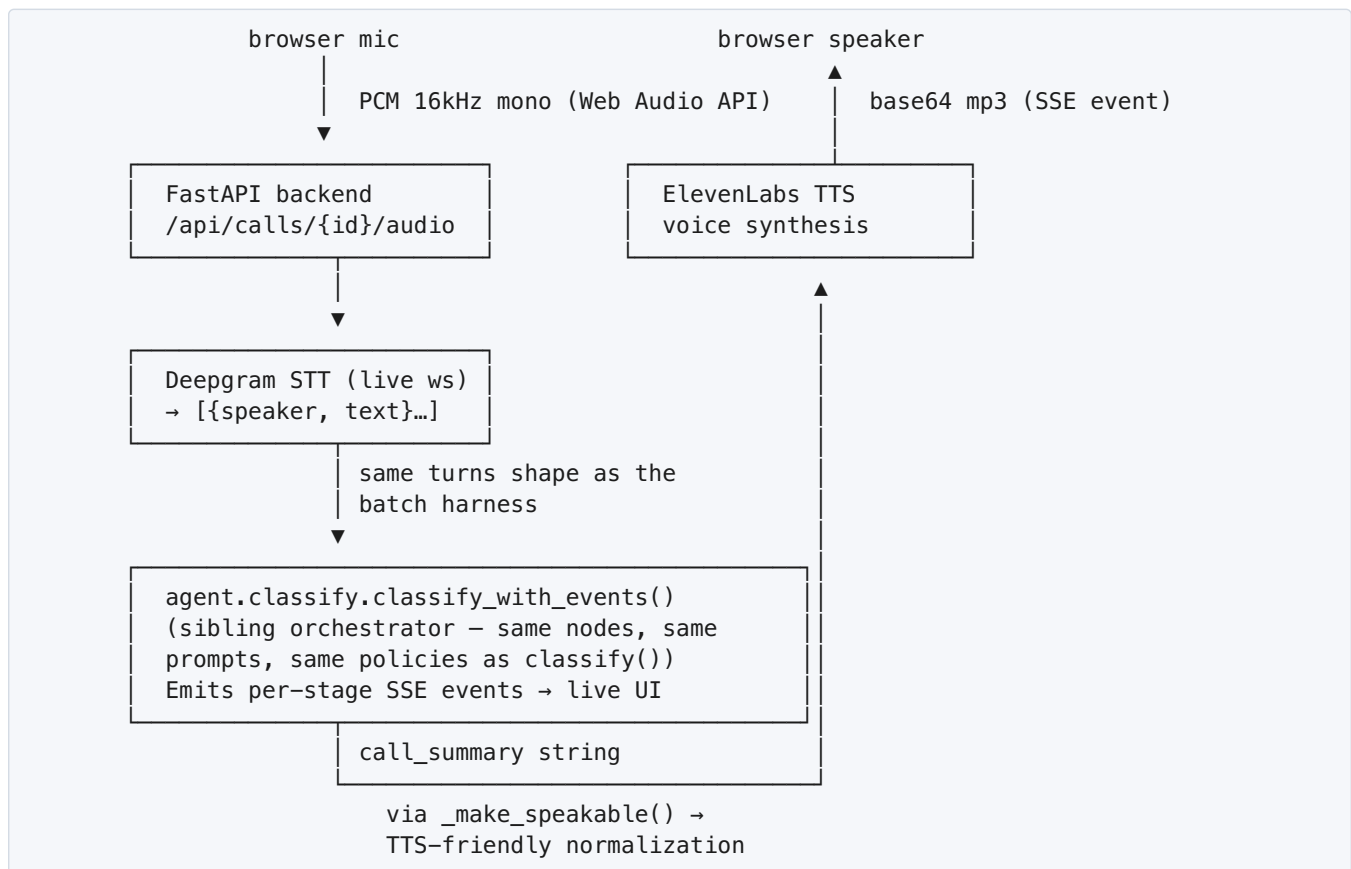
Each scoring axis in `evaluation/scoring.py::AXIS_WEIGHTS` maps to at least one section in this doc. Hard-fail penalty has a dedicated mechanism in §4.5.

Axis	Weight	Section(s)
Category accuracy	10%	§2 (classify node), §5 (RAG), §3.5 (boundary rules)
Subcategory accuracy	15%	§5.4 (audit-corrected retrieval), §3.5 (boundary rules)
Risk-level accuracy	10%	§3.1 (base risk), §3.2 (modifiers)
HITL trigger F1	15%	§3.3, §4
Clarification F1	5%	§7
Field-extraction accuracy	10%	§2 (extract + reconcile_location nodes)
Vendor-match accuracy	10%	§6
Auto-resolution rate	10%	§3.3 (rule precision), §7.3 (conservative clarify trigger)
Call-summary present	5%	§2 (summary node, PR #54)
Trainer-log present	10%	§8
False-911 penalty (–5/case)	—	§4.5 (life-safety whitelist + over-escalation guard)

12. Voice shell (bonus deliverable)

The brief lists a live STT → agent → TTS voice demo as a bonus axis. We built one. It is **not** on the grading path — the batch harness never touches it — but it exercises the same pipeline end-to-end with real audio and demonstrates that the contract holds at the demo layer too.

12.1 Architecture



Source: `backend/routes/calls.py`, `voice/session.py`, `frontend/components/VoiceCapture.tsx`.
Shipped behind PR #88 (peer-reviewed, pending merge at time of submission).

12.2 What's interesting

- **Multi-turn clarification.** If the agent's clarification node fires on the first pass, the backend speaks the question, re-opens the caller's microphone, captures the follow-up answer, appends those turns to the bus, and re-runs the pipeline against the full transcript. The summary stage is suppressed on the first pass when clarification fires, so callers don't hear a half-baked summary before they've finished talking.
- **Echo elimination.** The first version of the voice loop opened the microphone the moment the call started, which captured the agent's own TTS greeting through the speakers as caller speech. Deepgram transcribed the greeting as a fake caller turn ("CBRE maintenance, this is Dana...") and the SECURITY-aligned guard mis-classified it. Fix: the frontend registers a `greetingDoneCallback` per phase; the

microphone only opens when the corresponding TTS audio's `onended` fires. Same callback pattern for the clarification follow-up.

- **TTS normalization.** `_make_speakable()` rewrites strings before TTS: `"Floor 7"` → `"floor number 7"`, `"MEDIUM"` → `"medium"`, `"v_002"` → spelled out. This is cosmetic but removes the worst of the speech-engine pronunciation failures.
- **Live pipeline UI.** Each stage emits an SSE event on `started` → `complete`; the frontend's `PipelineDiagram` animates them sequentially. An `await asyncio.sleep(0)` after every publish in `backend/events/emitter.py` forces per-event SSE flushing — without it, all nine stages turned green at once when the pipeline finished.

12.3 Why a sibling orchestrator instead of wrapping `classify()`

`classify_with_events()` is a parallel orchestrator in `agent/classify.py`, **not** a wrapper around `classify()`. The output dict is byte-identical; the difference is two-fold:

1. Per-stage event emission for the live UI.
2. HITL pause/resume hooks for the reviewer modal — `classify()` runs straight through and is auto-approved by the batch harness; the voice loop's `pause_at_validator=True` actually waits for a reviewer click.

We considered factoring out a shared inner function and having both orchestrators wrap it. Rejected because (a) the byte-identical hash guard on `classify()` (`tests/test_classify_unchanged.py`) is load-bearing for the grader, and (b) the orchestration shells are short and the duplication is easy to keep in sync.

The shared assets are the node modules, the prompts, and the policy tables — those are the only places that could silently diverge, and they don't.

12.4 Zero impact on grading

The voice shell could be deleted tomorrow and the composite score would not move. The batch harness imports `agent.classify:classify` and runs through the synchronous orchestrator; nothing under `backend/`, `voice/`, or `frontend/` is on that path.

13. Limitations & known weaknesses

We list these because a graded design doc should be honest about what we know we don't have.

1. **Clarification triggers are heuristic, not learned.** §7.1 lists four rules (sparse caller, generic opening, floor self-correction, no location) tuned by hand against the dev set. A learned trigger over `(transcript, true_needs_clarification)` would generalise better, but we don't have post-clarification labels in the corpus.
2. **Vendor tie-breaker has no past-performance signal.** §6.3 picks by rating → cost → SLA → alphabetical. Real CBRE dispatch would weight by *this vendor's* historical success rate on *this subcategory*, which we don't compute. The data is present in `historical_records.json`; joining it into the tie-breaker is a real follow-up.
3. **LLM determinism is best-effort.** OpenAI's `seed` parameter is documented as best-effort; identical input occasionally yields different output across runs. We measure this directly: the baseline-recheck run (§9.7 row f-prime) drifted -0.44 composite from the corrected baseline on byte-identical code. Any change smaller than ± 0.45 is invisible at single-run resolution.
4. **No human evaluation of `call_summary`.** The 5% summary axis is credited for *presence*, not *quality*. We hit 100% on presence; the actual narrative quality is unmeasured. A separate LLM-as-judge pass over the 200 dev summaries would be the next step.
5. **Trainer-log replay is designed, not implemented.** §8.3 describes four replay paths into fine-tuning. We have not run a fine-tune. The format is correct; the loop is unbuilt.
6. **Voice latency budget is generous.** The voice shell adds ~2s for STT + ~3s for TTS on top of the ~5–10s pipeline. Total caller-perceived latency is ~10–15s per turn. Production would want streaming TTS (start speaking before the full string is generated) and faster STT end-pointing.
7. **RAG store is built on first run, not shipped.** We ship a build step in the README (~1–3 min, ~\$0.02 in embedding calls) rather than the `chroma_store/` artifact. If the grader's environment has no OpenAI embedding access, first-run build fails. The spec explicitly permits either approach; we picked the lighter option.
8. **Single-LLM-provider concentration.** Everything runs against OpenAI (`gpt-4o-mini` + `text-embedding-3-small`). A provider outage takes the whole agent down. Anthropic / local-LLM fallback would be a real production hardening.
9. **Per-axis sensitivity to noise floor.** The HITL-F1 axis at 0.748 could be 0.73 or 0.77 on the next run, purely from LLM drift. We rely on the 2.5σ headroom from §9.7 as evidence that levers 1+2 are real, but a single-run regression on a noise day could still look bad.

14. Reproducibility one-pager

The full README is at the repo root; this section is the cheat-sheet.

One-time setup.

```
git clone https://github.com/acm-industry/rohan_sachit_CBRE.git
cd rohan_sachit_CBRE
make install                                # creates .venv, installs deps
echo "OPENAI_API_KEY=sk-..." > .env      # required
python -m agent.rag.build_index             # ~1-3 min, ~$0.02
```

Grade against the dev set (200 transcripts, has labels).

```
python evaluation/run_eval.py \
  --agent your_submission.agent:classify \
  --eval evaluation/eval_transcripts_dev.json \
  --out eval_runs/dev_run.json

python evaluation/scoring.py \
  --eval evaluation/eval_transcripts_dev.json \
  --ground-truth evaluation/dev_labels.json \
  --predictions eval_runs/dev_run.json
```

Expected: composite $\approx 92.33 \pm 0.45$ (noise floor; see §9.7 row f-prime for the noise-measurement run).

Grade against the test set (800 transcripts, no labels). Same two commands with `--eval evaluation/eval_transcripts_test.json` and a held-out answer key. Shipped predictions file at repo-root `predictions.json` (tag `submission-v1`) was generated against agent SHA `c6b9255`; if grading against the final code, regenerate first.

Critical environment variables.

Variable	Required	Default	Purpose
OPENAI_API_KEY	yes	—	Chat + embedding calls
AGENT_CHAT_MODEL	no	gpt-4o-mini	Pinned per-submission
AGENT_CHAT_TEMPERATURE	no	0.0	Pinned for reproducibility
AGENT_CHAT_SEED	no	7	OpenAI seed (best-effort)
AGENT_EMBEDDING_MODEL	no	text-embedding-3-small	Changing requires <code>--force</code> rebuild of the chroma store
DEEPGRAM_API_KEY	voice only	—	STT for the voice demo (not graded)
ELEVENLABS_API_KEY	voice only	—	TTS for the voice demo (not graded)

Per-call latency. Hard timeout 30s in `evaluation/run_eval.py:46`. 1K-call dev run completes in ~17 minutes on a typical laptop; 800-call test-set run in ~70 minutes.

Hash guard. `tests/test_classify_unchanged.py` AST-hashes `classify()` and fails if the body drifts from the pinned SHA. Run via `make test-contract` (CI gate, < 1s).

15. Decision log

The big design decisions, with the rejected alternative and why:

#	Decision	Rejected alternative	Why
1	LangGraph for orchestration	Plain Python function-chain	We need explicit pause/resume semantics for HITL — <code>interrupt()</code> + <code>Command(resume=...)</code> + <code>update_state()</code> are first-class in LangGraph. Function chain would require hand-rolling the same primitives.
2	OpenAI <code>gpt-4o-mini</code>	Anthropic Claude / local Llama	OpenAI is the only provider with seed for reproducibility; <code>gpt-4o-mini</code> is cheap enough to run 1K-call evals iteratively (~\$0.10/run).
3	Per-subcategory base-risk from history, not taxonomy	Hand-curate from <code>operational/taxonomy.md</code>	The QA audit shows intake operators systematically over-state severity. Final (technician-on-site) labels are the same source the scorer uses.
4	Risk-weighted HITL gate (PR #89)	Hard low-confidence threshold for every call	Brief says $\text{Risk} = P(\text{error}) \times \text{Cost}(\text{error})$. A low-confidence LOW-risk call has trivial cost; pausing it just inflates the reviewer queue and depresses auto-resolution.
5	Audit-corrected RAG retrieval (§5.4)	Copy <code>intake_*</code> labels straight into the prompt	Re-trains the classifier on labels QA explicitly flagged as wrong. Demoting reclassified records and inlining <code>[QA-RECLASSIFIED from ...]</code> shows the LLM both the corrected label and the pattern of correction.
6	At-capacity vendors stay in candidate pool (PR #84 lever 1)	Hard-filter at-capacity vendors on emergencies	Brief says the cache is “intentionally stale”. Hard-filtering on a stale <code>at_capacity</code> was escalating 16 emergencies to nothing; dispatching the at-capacity vendor is strictly better when seconds count.
7	Demote <code>pipe_leak</code> / <code>power_outage</code> modal risk	Keep the modal-rule choice	Both subcategories are ~50/50 MEDIUM/HIGH

#	Decision	Rejected alternative	Why
	HIGH → MEDIUM (PR #84 lever 2)		historically; modal-rule was a coin flip that landed on HIGH and triggered HITL on every call. Soft/hard cue modifiers still escalate when language warrants.
8	Sibling orchestrator <code>classify_with_events()</code> for voice	Wrap <code>classify()</code> and add an event-emitter shim	The byte-identical hash guard on <code>classify()</code> is load-bearing for graders; a wrapper would either break the hash or require careful boundary management. Sibling orchestrator keeps the grading path frozen.
9	Build chroma index on first run, don't ship it	Commit the built <code>chroma_store/</code> artifact	Spec explicitly permits either; build-step keeps repo size down and lets graders use their preferred embedding model. Trade-off: requires OpenAI embedding access at grade time.
10	Reject PR #90 (prompt disambiguation hints)	Merge for putative classifier wins	Dev eval came back 92.27 (−0.06), with a NEW <code>waste_odor</code> → <code>fire_smoke</code> confusion that regressed <code>over_escalation_trap</code> 16/16 → 15/16. The signal was within noise floor and the regression was real; closed unmerged.

16. Glossary

For readers without AI/ML background. Terms appear in the order most useful for understanding this document.

- **LLM (Large Language Model).** A neural network trained on text that takes a string in and returns a string out. We use OpenAI's `gpt-4o-mini`. Treat it as a fuzzy function with no memory across calls.
- **Prompt.** The string you send to an LLM. Includes both the instructions (“classify this call into one of these subcategories...”) and the data (“transcript: ...”). Most of the engineering in this project is in the prompt for `classify`.
- **Token.** A unit of text the LLM operates on, roughly $\frac{3}{4}$ of a word. Pricing is per token. Our prompt is ~2K input tokens per call.
- **Temperature.** A knob from 0 to 1 controlling randomness in the LLM’s output. We pin `temperature=0` for determinism — same input → same output (within provider noise).
- **Seed.** A second determinism knob OpenAI exposes. With `temperature=0` and `seed=7`, the LLM is *mostly* deterministic; the remaining drift is what we call the “noise floor” (~±0.45 composite).
- **Structured output / function calling.** A mode where the LLM is forced to return JSON matching a schema you define, rather than free text. We use this for `extract` and `classify` so we get a guaranteed- parseable result.
- **Embedding.** A list of ~1500 numbers that represents the meaning of a piece of text. Two pieces of text with similar meaning have embeddings that are mathematically close. We embed all 10K historical records once with `text-embedding-3-small`.
- **Vector search / vector store.** A database that stores embeddings and lets you find the k most-similar ones to a query embedding. We use Chroma; the store lives at `agent/rag/chroma_store/`.
- **RAG (Retrieval-Augmented Generation).** The pattern of (1) embed the query, (2) vector-search for similar examples, (3) include those examples in the LLM’s prompt. Lets the LLM “remember” facts it wasn’t trained on. Our classifier is RAG-augmented over historical records.
- **Agent.** A pipeline that orchestrates multiple LLM calls (and other tools) to accomplish a task. “Agentic” just means multi-step. Our agent has 9 stages, of which 2 actually call an LLM.
- **LangGraph.** A Python library for building agents as directed graphs. Each node is a function; edges control flow. We use it for the built-in `interrupt()` and `Command(resume=...)` primitives.
- **interrupt().** A LangGraph call that pauses the agent mid-graph, hands a payload to a human reviewer, and waits. We use it at the validator gate (§4.1).
- **Command(resume=...).** The reviewer’s reply to an `interrupt()`. Tells the agent to continue with the reviewer’s decision baked in (approve/override).
- **HITL (Human in the Loop).** Any system that pauses for human approval before acting. Our HITL is at the validator gate; the trigger rule is in §3.3.
- **F1 score.** A combined precision + recall metric, 0 to 1. Our HITL-F1 of 0.748 means “if we say a call needs review, we’re right ~75% of the time, and we catch ~75% of the calls that actually need review.” It’s the harmonic mean of those two.
- **STT (Speech to Text).** Converts audio to text. We use Deepgram.
- **TTS (Text to Speech).** Converts text to audio. We use ElevenLabs.

17. Acknowledgments

- **Rohan Iyer** — co-author. Owned the LangGraph orchestration, the derived-policy pipeline (§3), the validator-gate primitives (§4), and the eval-mode auto-approve plumbing (§4.6). Drove the AC-AI deep analysis chats that surfaced levers 1+2 (§9.7 row g).
- **Sachit Madaan** — co-author. Owned the voice shell (§12: backend SSE, Deepgram/ElevenLabs integration, multi-turn clarification, echo elimination), the demo frontend, the spec walkthrough, and the design doc (this artifact). Drove the master-chat synthesis across all side-sessions.
- **CBRE × UCSB × ACM × Turing** — for the problem statement, the 10K-record historical corpus, the QA audit findings, the 1K-transcript dev/test split, and the scoring rubric.
- **Side-session contributors** — peer-driven PRs include #73 (security hardening), #77 (latency profiling), #84 (levers 1+2), #87 (presentation deck), #89 (risk-weighted HITL), and #91 (design-doc v2). Code review on #78, #88 (voice integration), and #92 (submission-layout copy) came from the same peers.

Last updated: 2026-05-22. Document version v3 (final submission). Authors: Rohan Iyer + Sachit Madaan. Current dev composite: 92.33 / 100, 0 false-911. Voice demo working end-to-end.